

nlp-assignment

September 2, 2024

SAMEER KHATWANI

PRN: 22070126099

```
[3]: import pandas as pd
import numpy as np
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
import re
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[4]: df = pd.read_csv('IMDB Dataset.csv')
df.head()
```

```
[4]:                                     review sentiment
0  One of the other reviewers has mentioned that ... positive
1  A wonderful little production. <br /><br />The... positive
2  I thought this was a wonderful way to spend ti... positive
3  Basically there's a family where a little boy ... negative
4  Petter Mattei's "Love in the Time of Money" is... positive
```

```
[8]: # Preprocessing
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('stopwords')

stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

def preprocess_text(text):
```

```

# Remove URLs
text = re.sub(r"http\S+|www\S+|https\S+", '', text, flags=re.MULTILINE)

# Remove punctuation and numbers
text = re.sub(r'\W', ' ', text)

# Tokenization
tokens = word_tokenize(text)

# Remove stopwords and lemmatize
tokens = [lemmatizer.lemmatize(word.lower()) for word in tokens if word.
↳lower() not in stop_words]

return ' '.join(tokens)

# Apply preprocessing to the 'review' column
df['review'] = df['review'].apply(preprocess_text)
df.head()

```

```

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\samee\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\samee\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\samee\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!

```

```

[8]:
      review sentiment
0  one reviewer mentioned watching 1 oz episode h... positive
1  wonderful little production br br filming tech... positive
2  thought wonderful way spend time hot summer we... positive
3  basically family little boy jake think zombie ... negative
4  petter mattei love time money visually stunnin... positive

```

```

[9]: # Vectorization
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['review'],
↳df['sentiment'], test_size=0.2, random_state=42)

# Vectorization
count_vect = CountVectorizer()
tfidf_vect = TfidfVectorizer()

X_train_count = count_vect.fit_transform(X_train)
X_test_count = count_vect.transform(X_test)

```

```
X_train_tfidf = tfidf_vect.fit_transform(X_train)
X_test_tfidf = tfidf_vect.transform(X_test)
```

```
[10]: # Step 6: Model Training and Evaluation
def evaluate_model(model, X_train, X_test, y_train, y_test):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # Print classification report
    print(classification_report(y_test, y_pred))

    # Plot confusion matrix
    cm = confusion_matrix(y_test, y_pred, labels=model.classes_)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=model.
    ↪classes_, yticklabels=model.classes_)
    plt.ylabel('Actual')
    plt.xlabel('Predicted')
    plt.show()

# Logistic Regression
print("Logistic Regression with CountVectorizer")
evaluate_model(LogisticRegression(), X_train_count, X_test_count, y_train,
    ↪y_test)

print("Logistic Regression with TfidfVectorizer")
evaluate_model(LogisticRegression(), X_train_tfidf, X_test_tfidf, y_train,
    ↪y_test)

# SVC
print("SVC with CountVectorizer")
evaluate_model(SVC(), X_train_count, X_test_count, y_train, y_test)

print("SVC with TfidfVectorizer")
evaluate_model(SVC(), X_train_tfidf, X_test_tfidf, y_train, y_test)

# Random Forest
print("Random Forest with CountVectorizer")
evaluate_model(RandomForestClassifier(), X_train_count, X_test_count, y_train,
    ↪y_test)

print("Random Forest with TfidfVectorizer")
evaluate_model(RandomForestClassifier(), X_train_tfidf, X_test_tfidf, y_train,
    ↪y_test)
```

Logistic Regression with CountVectorizer

C:\Users\samee\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n

```
2kfra8p0\LocalCache\local-packages\Python312\site-  
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:


https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
    precision    recall  f1-score   support  
  
   negative      0.89     0.88     0.88     4961  
   positive      0.88     0.89     0.89     5039  
  
   accuracy              0.88     10000  
  macro avg      0.88     0.88     0.88     10000  
 weighted avg      0.88     0.88     0.88     10000
```




Logistic Regression with TfidfVectorizer

	precision	recall	f1-score	support
negative	0.91	0.88	0.89	4961
positive	0.88	0.91	0.90	5039
accuracy			0.90	10000
macro avg	0.90	0.90	0.90	10000
weighted avg	0.90	0.90	0.90	10000




SVC with CountVectorizer

	precision	recall	f1-score	support
negative	0.90	0.86	0.88	4961
positive	0.87	0.91	0.89	5039
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000




SVC with TfidfVectorizer

	precision	recall	f1-score	support
negative	0.91	0.89	0.90	4961
positive	0.89	0.92	0.90	5039
accuracy			0.90	10000
macro avg	0.90	0.90	0.90	10000
weighted avg	0.90	0.90	0.90	10000



Random Forest with CountVectorizer				
	precision	recall	f1-score	support
negative	0.86	0.86	0.86	4961
positive	0.86	0.86	0.86	5039
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000



Random Forest with TfidfVectorizer				
	precision	recall	f1-score	support
negative	0.85	0.87	0.86	4961
positive	0.86	0.85	0.86	5039
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

